


AVL Balanced Trees, Rotations & Splay Trees

Paul Hayter

Trees


- Tree terminology



Nyhoff, ADTs, Data Structures and Problem Solving with C++, Second Edition, © 2005 Pearson Education, Inc. All rights reserved. 0-13-140909-3

- **Depth** – how many levels to the (sub)root is a node (looking up)
- **Height** – how many levels from a (sub)root to deepest child (looking down)
- **Height** is independent of the **subtree** we are considering, because **height** is measured *downward*, not upward. Height is a key feature of AVL trees.

Balanced and Unbalanced Trees




For both trees, what is the height of the two subtrees of the root?

An **AVL balanced tree** is when subtree heights differ by no more than 1.
Which tree is balanced? Why?

Balance Factor = Height(Right subtree) – Height(Left subtree)
(though some texts write this the other way)

AVL Balanced Trees



This tree meets **AVL condition** at the root (**5**), since its left subtree (**2**) has height 2, and its right subtree (**8**) has height 1.


Balance Factor = $1 - 2 = -1$

This tree **violates AVL condition** at the root node (**7**) because its left subtree (**2**) has height 2, while its right subtree (**8**) has height 0.

Balance Factor = $0 - 2 = -2$.

Except for the root, *both* trees have **AVL** balanced nodes, so in this case, it is only the root node on the right example that is not **AVL** balanced.

Computing Height




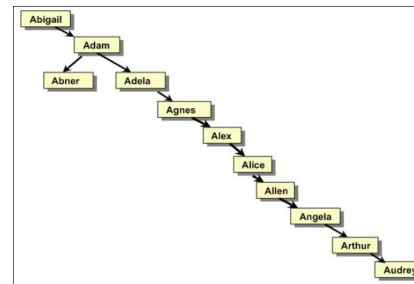
$$\text{Height} = \max(\text{leftChildHeight}, \text{rightChildHeight}) + 1$$

Discussion:

- What are the heights of the root nodes of these trees?
- What is the height of a null node (i.e., a child of a leaf node)?
(Hint: consider the children of Node 8 of the unbalanced tree on right.)

Balanced Trees

- Discussion:
 - Is an AVL tree a Binary Search Tree? Explain.
 - What is the balance condition for an AVL Tree?
- Worked Exercise
 - Draw an Unbalanced Tree
 - Draw a Balanced Tree
- What advantage does a Balanced Tree have over an unbalanced Tree?
- What is the Big-O of searching a Balanced Tree?



Rotations

Discussion:


- What is a rotation and when is it used?
- Weiss' text notation (which is common) is:
 - Right rotation when nodes move to the right
 - Left rotation when nodes move to the left
 - Note: Loceff modules have the opposite meaning

Note: Be aware that double rotation terminology in the Weiss text is different from the rotation language used in the Loceff modules.

Loceff Module	Weiss' Text
Double Left	Right-Left
Double Right	Left-Right


Rotations: Right Rotation Example

Initial unbalanced tree



Weiss Notation (opposite of Loceff) for 'left' & 'right'

Discussion: Why is this balance operation legal for a Binary Search Tree (BST)?



Note that Y is still less than k_2 and greater than k_1



Final balanced tree


Illegal tree: intermediate state

For visualization only!!!

Note:

$k_1 \rightarrow X$ unchanged

$k_2 \rightarrow Z$ unchanged




Rotations: Left Rotation Example

Weiss Notation (opposite of LocEFF) for 'left' & 'right'


Initial unbalanced tree

Final balanced tree




Rotation Exercises

Weiss Notation (opposite of LocEFF) for 'left' & 'right'



Problem 1




Problem 2


Note: For practice only – these rotations won't balance the trees.
Also, Node letters are for ID only and are not keys.

Rotation Exercises

Weiss Notation (opposite of LocEFF) for 'left' & 'right'




Problem 3



Problem 4

Note: For practice only – these rotations won't balance the trees.
Also, Node letters are for ID only and are not keys.

Rotation Exercises




Discuss:

- Is this tree balanced?
- Is it balanced after insertion into Y-subtree with height increase?
- Is it balanced after insertion into Z-subtree with height increase?


Remember: $\text{Height} = \max(\text{left_child_height}, \text{right_child_height}) + 1$

Rotation Exercises

Weiss Notation (opposite of Loeff) for 'left' & 'right'



White Tree



Grey Tree

1. Rotate Left both trees and annotate changes in height (the h values)
2. After rotation, which tree(s) are balanced?
3. How could a balanced outcome be predicted?
4. Is this prediction the same for a Rotate Right?


Remember: Height = $\max(\text{leftChildHeight}, \text{rightChildHeight}) + 1$

Double Rotation

In Left-Right Case rotate into a Left-Left Case to set up for a balanced rotation


Left-Right Case → Left-Left Case → Balance

Note how this is only necessary when the tree is 'heavy' on the **inside**, that is, of greater height on the tree's inside



Note the symmetry in applying the 'heavy inside' rule in both cases

Double Rotation Exercise



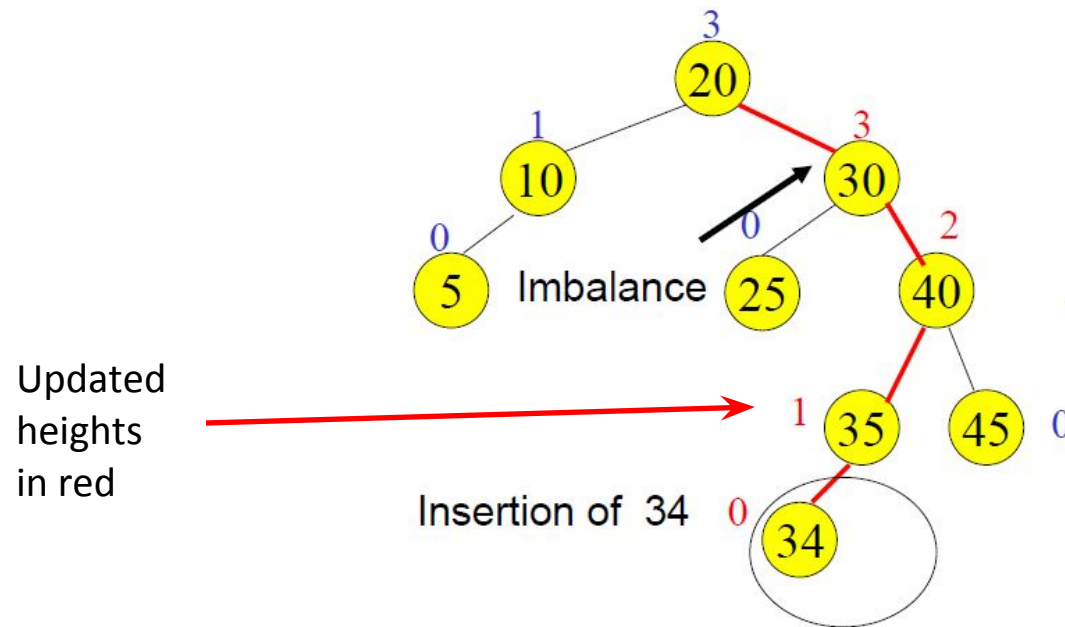
History of this tree

- 1) Inserted Node 34
- 2) Updated Node heights going up from node 34 one parent at a time (updated heights shown in red)
- 3) Checked Balance Factor after each height update

Based on Image from:

courses.cs.washington.edu/courses/cse373/06sp/handouts/lecture12.pdf

Double Rotation Exercise

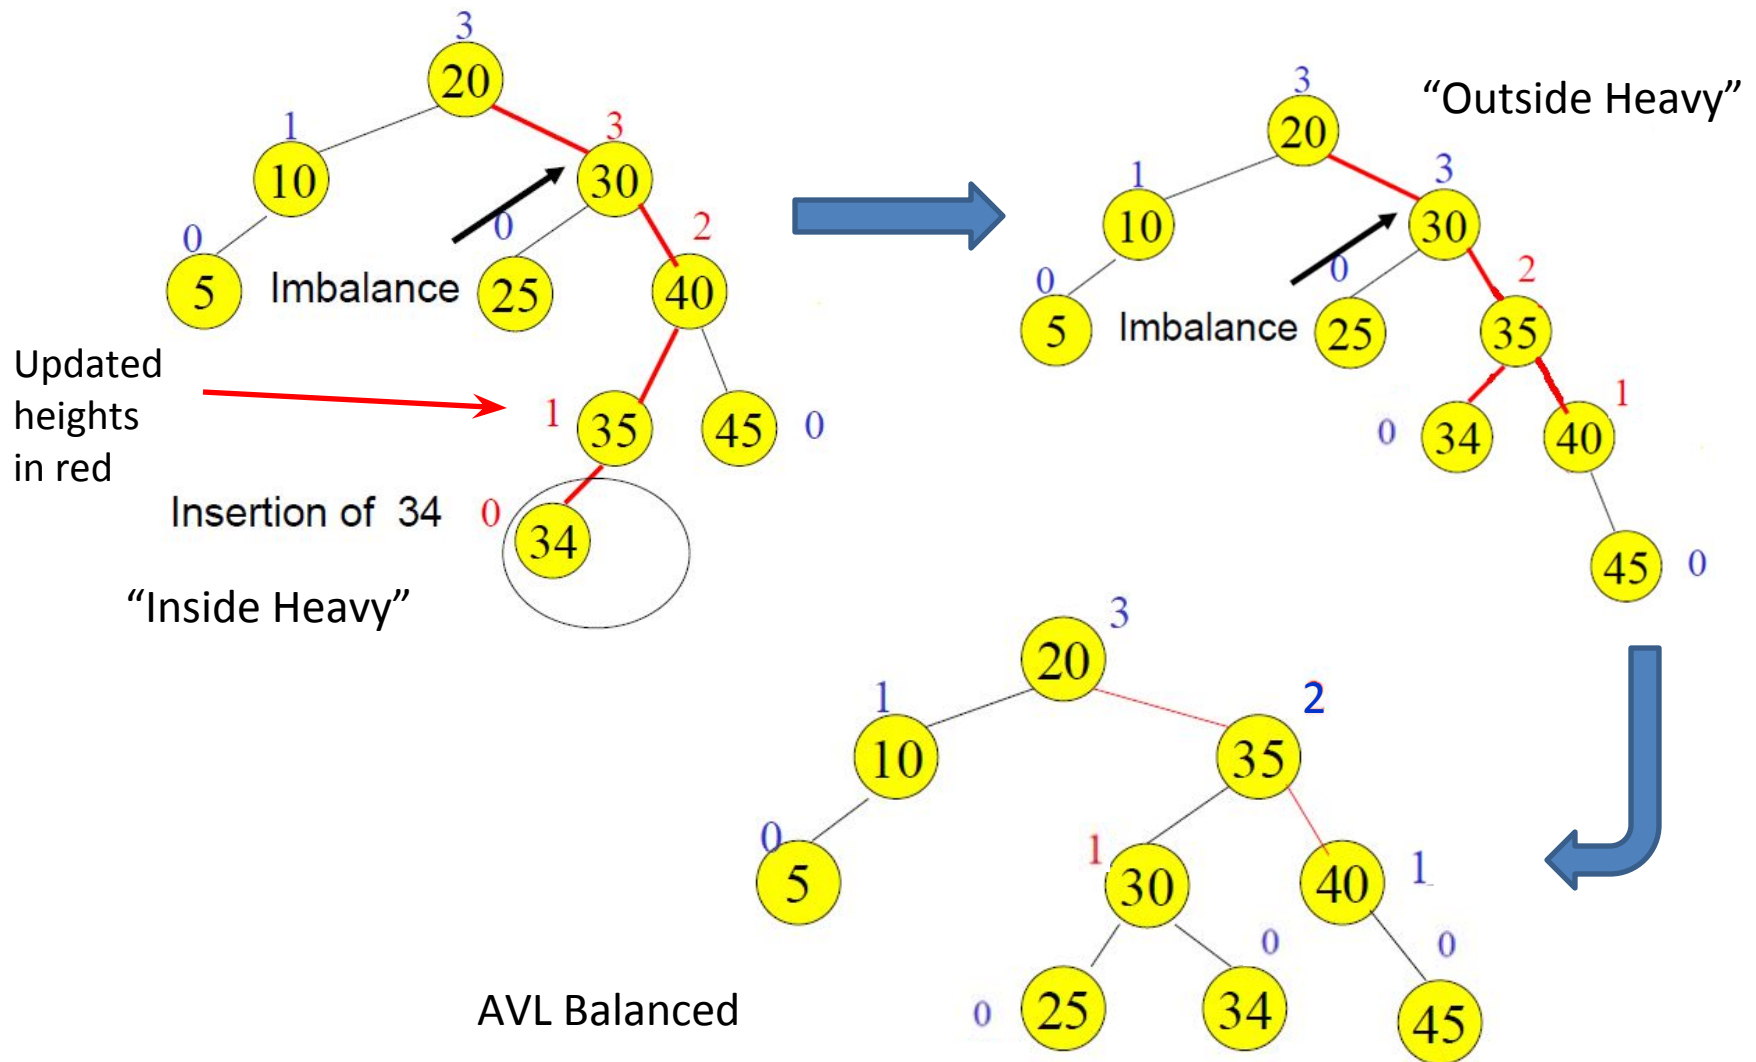


- 1) Why imbalance at Node 30?
- 2) Is tree 'heavy' on the inside?
- 3) What rotations should be done and at which nodes?
- 4) What is the result of these rotations?

Images:

<https://courses.cs.washington.edu/courses/cse373/06sp/handouts/lecture12.pdf>

Double Rotation Exercise



AVL Rotation Summary

Let the node that needs rebalancing be α .

There are 4 cases:

Outside Cases (require single rotation) :





1. Insertion into **left** subtree **of left** child of α .
2. Insertion into **right** subtree **of right** child of α .

Inside Cases (require double rotation) :


3. Insertion into **right** subtree **of left** child of α .
4. Insertion into **left** subtree **of right** child of α .

The rebalancing is performed through four separate rotation algorithms.

Note: Arrows show first (or only) rotation.




AVL Tree?



Is this an AVL Tree? Why or why not?

Splay Trees

- Discussion: Is a Splay Tree a Balanced Tree? Explain.
- What are the advantages of a Splay Tree?
- How might those advantages affect when this choice of tree is used?








Example Splay Tree from Weiss
Text

Splay Algorithm Explained

- First Pass Explanation
 - Top-level and general description
 - Rough sketches missing details
- Second Pass Explanation with navigation details
 - More detailed description
 - More detailed sketches

Splay Tree Algorithm Concept

Search for X

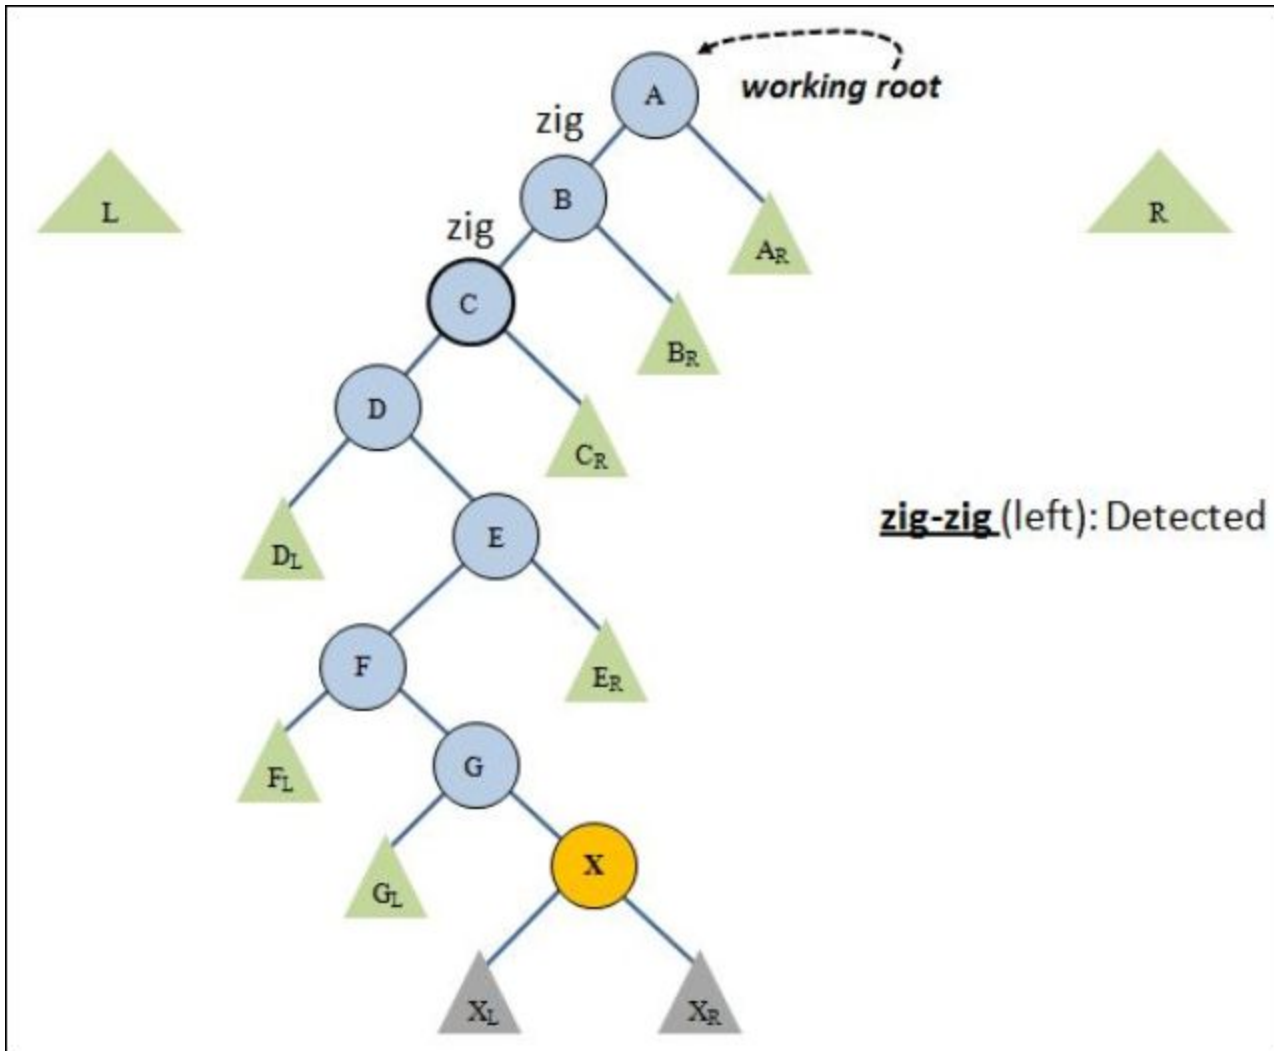


From Module 5B.4.3 by Michael LocEFF

Splay Tree Navigation (Detail)

- As we search down the tree looking for X, the following can happen:
 - We move twice to the Right or Left (through those children) – called a *zig-zig*
 - We move Right-Left or Left-Right (through those children – called a *zig-zag*
- Need to know how to disposition those nodes that we move through

Splay Tree Navigation



A *zig-zig* is TWO sequential data “checks” in the same direction when looking for X. Left-Left or Right-Right

Now we know:


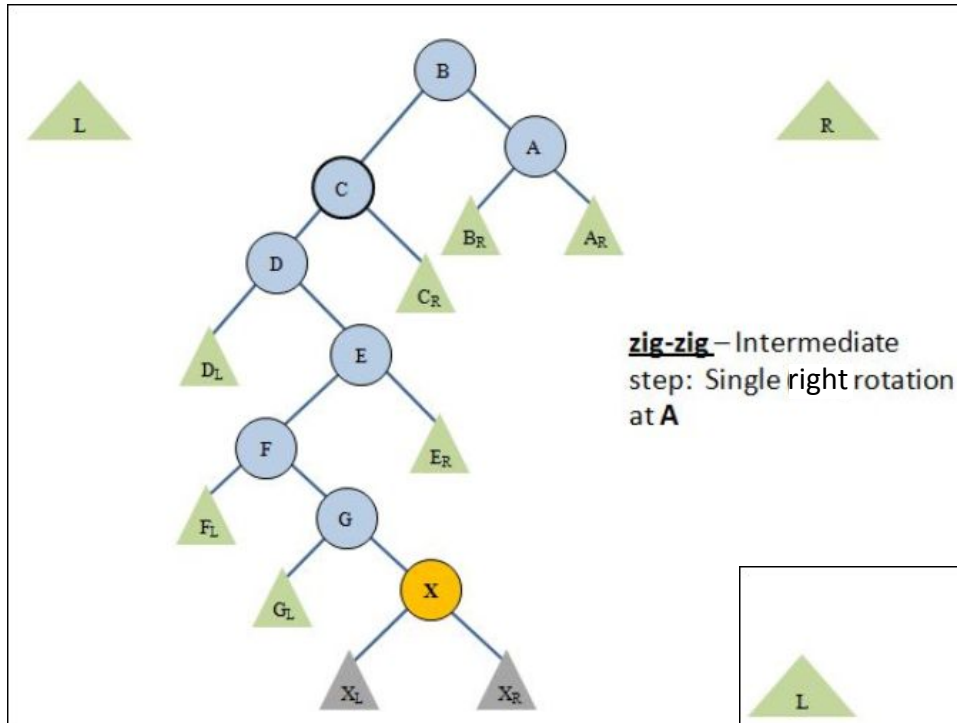
- 1) X is not A or B
- 2) Since Left-Left, then $X < A \ \& \ B$

Next Steps:

- 1) Single rotation Right at A
- 2) Remove A & B to minimum of R-Tree
- 3) Appoint C as new root

Note: this image starts at an *intermediate* phase of splaying so L and R have content

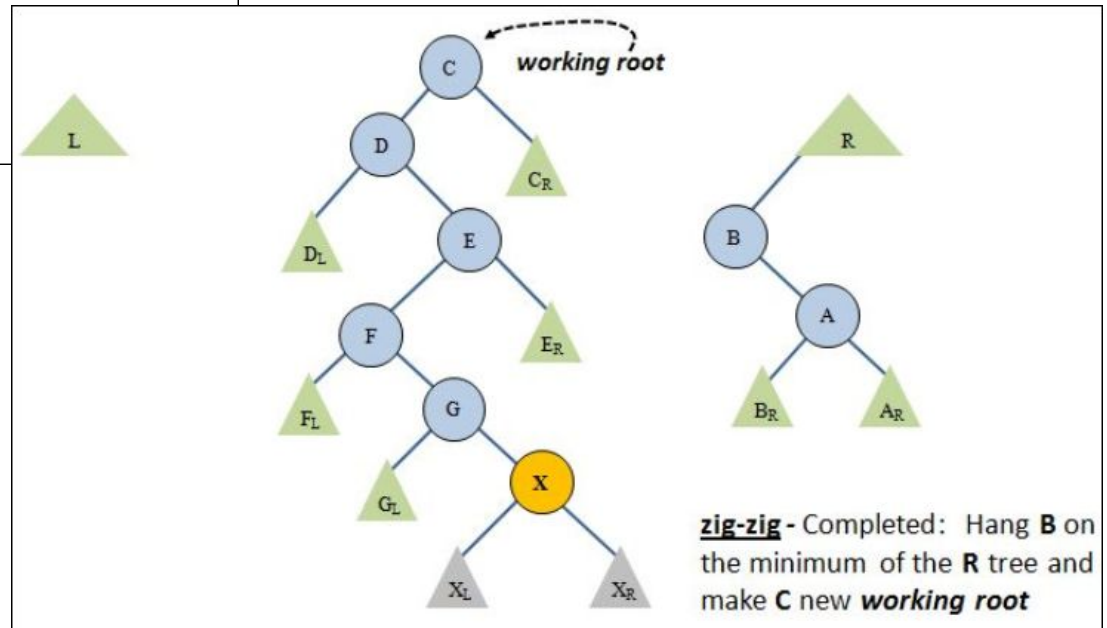
Splay Tree Navigation: zig-zig





Remember:
Moved to R-Tree since $X < A$ & B

If *zig-zig* had been Right-Right, these operations would be reversed since $X > A$ & B:

- 1) Single Left rotation
- 2) Move to L-Tree maximum



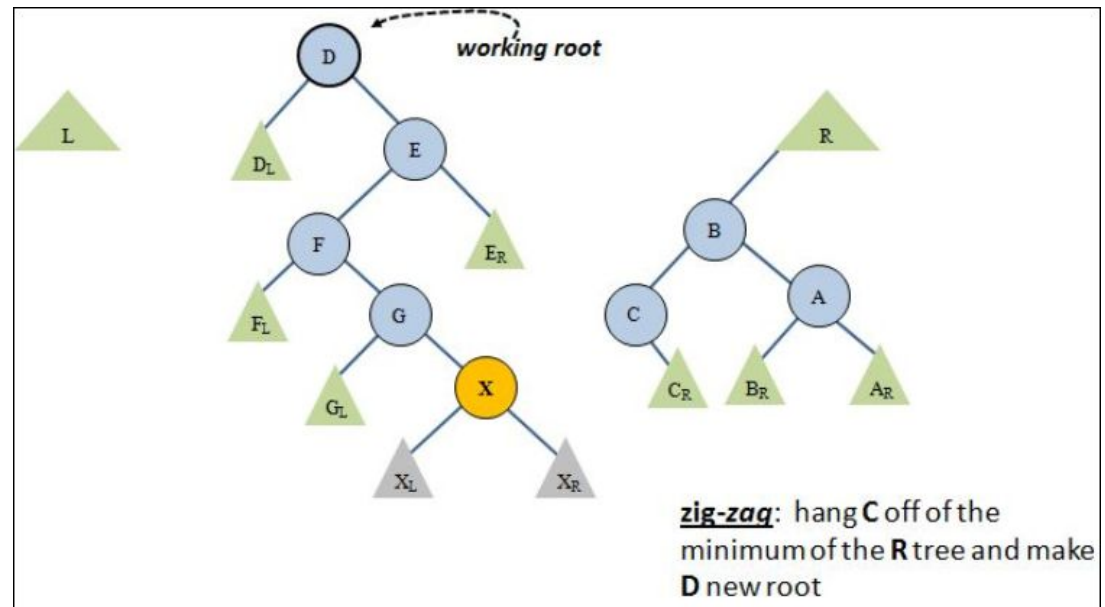
Splay Tree Navigation: zig-zag



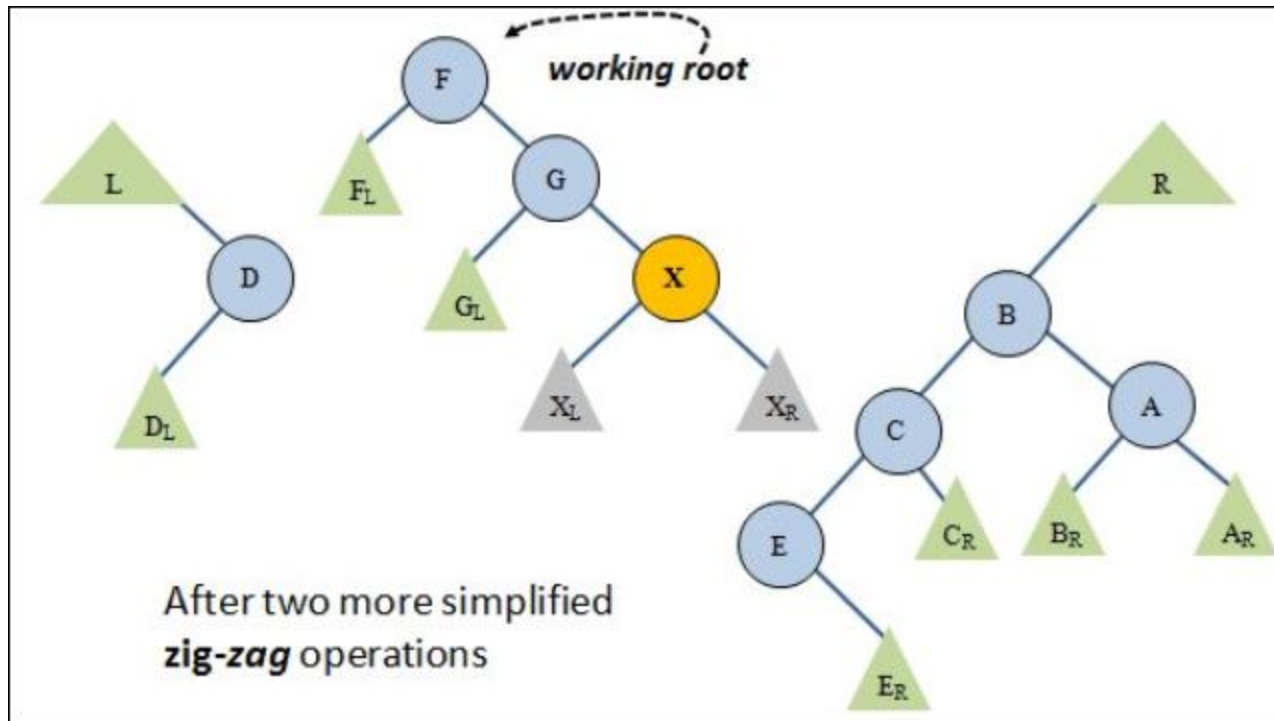
Remember:
Moved to R-Tree since $X < C$

If *zig-zag* had been Right-Left, these operations would be reversed since $X > C$:

- 1) Move to L-Tree maximum
- 2) (No rotation in either)




Splay Tree Navigation: Discussion



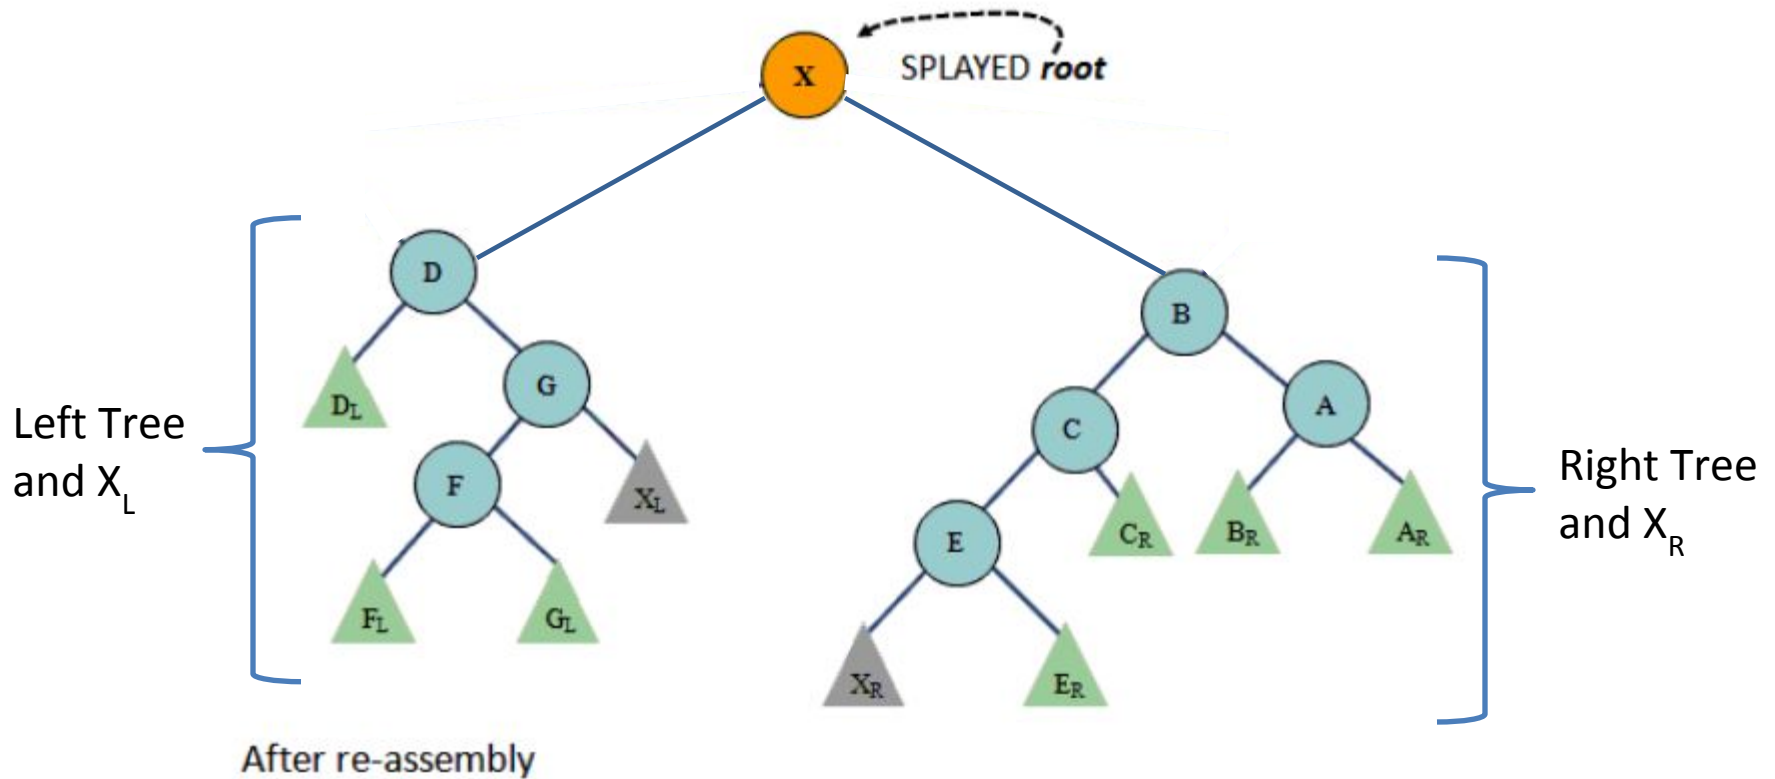
Discussion: What operations should be done now?

- 1) Rotate Left at F
- 2) Move F & G to L tree to the maximum position (since $F \& G < X$)

Splay Tree Navigation: X Found



Splay Tree Navigation: Re-assembly




Note location of X_L and X_R
Why does that move make sense?

Discussion: Splay Tree Observations

- Is the resultant Splay Tree balanced?
- How does the structure of the resultant Splay Tree compare with how it started?
- If there is a search for X again, how long will it take? What is the Big-O in that case?


Worked Exercise

Splay for 19 in this tree



Discussion:

- Knowing that 19 is not in this tree, how useful will the resultant Splay tree be if we first searched for 19, then wanted to insert 19?
- What would be the time complexity of the insertion after the search for 19?



Solution for Splay for 19 From Weiss' text

Note: Weiss algorithm is slightly different than the modules algorithm and has a "simplified zig-zag" in addition to "zig-zag".

Homework Advice

- Follow the pseudo-code in modules closely (sections 5B.6.1 through 5B.6.3)
 - Like the coding the Subset Sum algorithm, not all details for efficient implementation are given in the pseudo-code.
 - Be careful in maintaining BST structure in the right-tree, left-tree and main-tree.
- Use the rotate methods from the modules (make the small necessary modifications)


Caution: Due to slight differences in the Module's algorithm from the text by Mark Weiss, the text's Splay tree examples may differ from your 'play computer' results.

Backup


Splay Algorithm

- It is an iterative (loop) process that always has three trees
 - The **main tree** that you are searching (which you expect contains x)
 - A **left tree, L**, which contains nodes **less than x** that you have already examined and can be safely moved "out of the way" of the main tree, and
 - A **right tree, R**, which contains nodes **greater than x** that you have already examined and can be safely moved "out of the way" of the main tree
- Initially the main tree is the entire tree, **T**, that you are searching and **L** and **R** are empty trees.
- As the splay progresses, you will be moving left and right down the path towards x , much like we have done in the past in such algorithms. As you do, you will be removing nodes you encounter from the **main tree** and placing them into either **L** or **R**, depending on whether they are less than or greater than x . Thus, the **main tree** will get smaller, and the **L** and **R** trees will grow.
- If you find x (or hit a **null**), you are done with the splay. You can then **re-assemble** the three trees by adjusting a few links and return.


Rotation Exercise Answers




Solution to #1



Solution to #2




Solution to #3



Solution to #4


Height of NULL




In this tree:

- Zebra's left-child subtree height is 0
- Zebra's right-child subtree height is -1
- the difference in subtree height between these two children is 1, therefore this is a balanced tree

Double Rotation




Problem 1




Problem 2

Double Rotation Solutions




Problem 1




Problem 2

White Tree/Grey Tree Solution



Tree is the same for both trees except for “White Tree” the height of Y is 15 and the height of Z is 16 (unchanged from initial condition).


X_r Goes to Right Tree in Re-assembly




Note in both cases, X_r is greater than X and also less than A & B , therefore at re-assembly X_r goes to Left-Child of Right Tree (which has A & B above it).
 A & B are above since a zig-zig left causes their rotation and move to Right Tree.

Backup & Extra


AVL Tree of Maximum Height



AVL Insert & Balance Exercise




Insert 21



After insert is tree AVL misbalanced? If so, where?

What rotations needed to restore balance? Draw on whiteboard.



Splay Pseudo-code

- initialize the **rightTree**, **leftTree**, **rightTreeMin**, and **leftTreeMax** to **NULL**
- loop while **root** **!=** **NULL** (**root** should not *become* **NULL**, but this protects against **NULL** parameter)
 - if **x** < **root**
 - check for **root**'s left child **NULL**. If so, **x** not in tree, break loop.
 - if **x** < **root**'s left child we have **zig zig** (left) so do a single rotate (left) at **root**
 - check for (new) **root**'s left child **NULL**. If so, **x** not in tree, break loop.
 - add **root** to **rightTree** at its minimum node - update the **rightTreeMin** to point to this node
 - update the new working **root**: set **root** to **root**'s left child
 - otherwise, if **x** > **root**
 - check for **root**'s right child **NULL**. If so, **x** not in tree, break loop.
 - if **x** > **root**'s right child we have **zig zig** (right) so do a single rotate (right) at **root**
 - check for (new) **root**'s right child **null**. If so, **x** not in tree, break loop.
 - add **root** to **leftTree** at its maximum node - update the **leftTreeMax** to point to this node
 - update the new working **root**: set **root** to **root**'s right child
 - otherwise we have found **x** at **root**. break.
- reassemble
 - if the left tree is not **NULL**, hang **root**'s left child onto its maximum and set **root**'s left child = the left tree.
 - if the right tree is not **NULL**, hang **root**'s right child onto its minimum and set **root**'s right child = the right tree.